

Artificial Intelligence Term Project (Horses or Humans)

Dr. Aziz Al-otaibi

Stuednt name :

Abdulmajeed Mohammed Al-amri - 43803419

Meshal Bandar Hadhrawi – 43807307

Dataset information:

- Horses or Humans is a dataset of 300×300 images, created by Laurence Moroney, that is licensed CC-By-2.0 for anybody to use in learning or testing computer vision algorithms.
- The set contains 500 rendered images of various species of horse in various poses in various locations. It also contains 527 rendered images of humans in various poses and locations
- number of training images = 1027
- image size 300×300 width x height.
- Number of classes 2 labels.

Load Dataset :

- Loading the dataset from google URL (zip file).
- Using OS library give us access to the zipfile to unzip the file or data.

```
[43] !wget --no-check-certificate \
      https://storage.googleapis.com/laurencemoroney-blog.appspot.com/horse-or-human.zip \
      -O /tmp/horse-or-human.zip

--2020-12-04 23:15:04-- https://storage.googleapis.com/laurencemoroney-blog.appspot.com/horse-or-human.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.7.176, 142.250.73.208, 142.250.73.240, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.7.176|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 149574867 (143M) [application/zip]
Saving to: '/tmp/horse-or-human.zip'

/tmp/horse-or-human 100%[=====>] 142.65M  295MB/s   in 0.5s

2020-12-04 23:15:05 (295 MB/s) - '/tmp/horse-or-human.zip' saved [149574867/149574867]
```

```
[13] import os
      import zipfile

      local_zip = '/tmp/horse-or-human.zip'
      zip_ref = zipfile.ZipFile(local_zip, 'r')
      zip_ref.extractall('/tmp/horse-or-human')
      zip_ref.close()
```

Display a batch of 8 horses and 8 humans pictures:

```
[28] %matplotlib inline
```

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
```

```
# Parameters for our graph; we'll output images in a 4x4 configuration
nrows = 4
ncols = 4
```

```
# Index for iterating over images
pic_index = 0
```




```
# Set up matplotlib fig, and size it to fit 4x4 pics
fig = plt.gcf()
fig.set_size_inches(ncols * 4, nrows * 4)
```

```
pic_index += 8
next_horse_pix = [os.path.join(train_horse_dir, fname)
                  for fname in train_horse_names[pic_index-8:pic_index]]
next_human_pix = [os.path.join(train_human_dir, fname)
                  for fname in train_human_names[pic_index-8:pic_index]]
```

```
for i, img_path in enumerate(next_horse_pix+next_human_pix):
    # Set up subplot; subplot indices start at 1
    sp = plt.subplot(nrows, ncols, i + 1)
    # Don't show axes (or gridlines)
    sp.axis('on')
```

```
img = mpimg.imread(img_path)
plt.imshow(img)
```

```
plt.show()
```



Model :

❖ Building a Small Model from Scratch :

1.import tensorflow.

2.add convolutional layers

```
import tensorflow as tf
model = tf.keras.models.Sequential([
    # Note the input shape is the desired size of the image 300x300 with 3 bytes color
    # This is the first convolution
    tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    # The second convolution
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    # The third convolution
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    # The fourth convolution
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    # The fifth convolution
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    # Flatten the results to feed into a DNN
    tf.keras.layers.Flatten(),
    # 512 neuron hidden layer
    tf.keras.layers.Dense(512, activation='relu'),
    # Only 1 output neuron. It will contain a value from 0-1 where 0 for 1 class ('horses') and 1 for the other ('humans')
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

❖ The model.summary() method call prints a summary of the NN

```
[ ] model.summary()
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_1 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_2 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 35, 35, 64)	0
conv2d_3 (Conv2D)	(None, 33, 33, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_4 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 64)	0
flatten (Flatten)	(None, 3136)	0
dense (Dense)	(None, 512)	1606144
dense_1 (Dense)	(None, 1)	513

Total params: 1,704,097
Trainable params: 1,704,097
Non-trainable params: 0

Training(Epoch,Batch_size,Activation_function)

Note:

- Batch_size : 128
- Activation_function : ReLU
- ReLU stands for rectified linear unit, and is a type of activation function.
- We have 1027 images in the dataset and 2 classes

```
▶ from tensorflow.keras.preprocessing.image import ImageDataGenerator

# All images will be rescaled by 1./255
train_datagen = ImageDataGenerator(rescale=1/255)

# Flow training images in batches of 128 using train_datagen generator
train_generator = train_datagen.flow_from_directory(
    '/tmp/horse-or-human/', # This is the source directory for training images
    target_size=(300, 300),
    batch_size=128,
    # Since we use binary_crossentropy loss, we need binary labels
    class_mode='binary')
```

Found 1027 images belonging to 2 classes.

Training(Epoch,Batch_size,Activation_function)

Note:

- Steps_per_Epoch
= (1027 / 128) = 8
- verbose = 1, which includes both progress bar and one line per epoch

```
▶ history = model.fit(  
    train_generator,  
    steps_per_epoch=8,  
    epochs=15,  
    verbose=1)
```

Epoch 1/15
8/8 [=====] - 5s 677ms/step - loss: 0.7410 - accuracy: 0.5829
Epoch 2/15
8/8 [=====] - 6s 764ms/step - loss: 0.7678 - accuracy: 0.6630
Epoch 3/15
8/8 [=====] - 5s 674ms/step - loss: 0.3936 - accuracy: 0.8610
Epoch 4/15
8/8 [=====] - 5s 671ms/step - loss: 0.2927 - accuracy: 0.8710
Epoch 5/15
8/8 [=====] - 6s 747ms/step - loss: 0.3728 - accuracy: 0.8574
Epoch 6/15
8/8 [=====] - 6s 753ms/step - loss: 0.1532 - accuracy: 0.9492
Epoch 7/15
8/8 [=====] - 5s 664ms/step - loss: 0.1652 - accuracy: 0.9388
Epoch 8/15
8/8 [=====] - 5s 670ms/step - loss: 0.0538 - accuracy: 0.9833
Epoch 9/15
8/8 [=====] - 5s 663ms/step - loss: 1.7270 - accuracy: 0.8287
Epoch 10/15
8/8 [=====] - 5s 665ms/step - loss: 0.2149 - accuracy: 0.9210
Epoch 11/15
8/8 [=====] - 5s 659ms/step - loss: 0.0695 - accuracy: 0.9711
Epoch 12/15
8/8 [=====] - 6s 754ms/step - loss: 0.0241 - accuracy: 0.9944
Epoch 13/15
8/8 [=====] - 5s 652ms/step - loss: 0.0287 - accuracy: 0.9933
Epoch 14/15
8/8 [=====] - 5s 668ms/step - loss: 0.0245 - accuracy: 0.9889
Epoch 15/15
8/8 [=====] - 5s 674ms/step - loss: 0.0272 - accuracy: 0.9889

Result:

- You can click on Choose File, after selecting the file the test will be performed and give the result



```
import numpy as np
from google.colab import files
from keras.preprocessing import image

uploaded = files.upload()

for fn in uploaded.keys():

    # predicting images
    path = '/content/' + fn
    img = image.load_img(path, target_size=(300, 300))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)

    images = np.vstack([x])
    classes = model.predict(images, batch_size=10)
    print(classes[0])
    if classes[0]>0.5:
        print(fn + " is a human")
    else:
        print(fn + " is a horse")
```

Choose Files No file chosen

Cancel upload

Result:

After selecting an image from the horses file

```
import numpy as np
from google.colab import files
from keras.preprocessing import image

uploaded = files.upload()

for fn in uploaded.keys():

    # predicting images
    path = '/content/' + fn
    img = image.load_img(path, target_size=(300, 300))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)

    images = np.vstack([x])
    classes = model.predict(images, batch_size=10)
    print(classes[0])
    if classes[0]>0.5:
        print(fn + " is a human")
    else:
        print(fn + " is a horse")
```

Choose Files horse01-0.png

- **horse01-0.png**(image/png) - 151700 bytes, last modified: 10/22/2019 - 100% done
Saving horse01-0.png to horse01-0.png
[0.]
horse01-0.png is a horse

After selecting an image from the human file

```
import numpy as np
from google.colab import files
from keras.preprocessing import image

uploaded = files.upload()

for fn in uploaded.keys():

    # predicting images
    path = '/content/' + fn
    img = image.load_img(path, target_size=(300, 300))
    x = image.img_to_array(img)
    x = np.expand_dims(x, axis=0)

    images = np.vstack([x])
    classes = model.predict(images, batch_size=10)
    print(classes[0])
    if classes[0]>0.5:
        print(fn + " is a human")
    else:
        print(fn + " is a horse")
```

Choose Files human01-00.png

- **human01-00.png**(image/png) - 160725 bytes, last modified: 10/22/2019 - 100% done
Saving human01-00.png to human01-00.png
[1.]
human01-00.png is a human

Conclusion

- Finally In this project, we selected a dataset (Human_Horses) that is in Tenserflow, and built a neural network that uses more than one layer, and train the model and achieving high accuracy with minimum lost , and give the data for the training model and predicted the results.

3 Dec 2020